

THE GEORGE  
WASHINGTON  
UNIVERSITY

WASHINGTON, DC

# DISTRIBUTED SYSTEMS CS6421

## INTRO TO DISTRIBUTED SYSTEMS AND THE CLOUD COMPUTING

Prof. Roozbeh Hagnazar

Slides Credit:

Prof. Tim Wood and Prof. Roozbeh Hagnazar

# PROF. TIM WOOD

- Research:  
Virtualization  
platform design, cloud  
resource management,  
and software-based  
networking
- Teaching: Distributed  
Systems, Networking,  
Software Engineering,  
Senior Design





# PROF. ROOZBEH HAGHAZAR

- Started Programming in 1991 with Commodore 64
- Played several roles in technology, such as Developer, Modeler, Designer, Architect, Leader, CTO, etc.
- Teach Software Eng., Distributed Systems, Data Base Design Principles, Data Visualization, Operating System.



# ABOUT THIS COURSE

- Be prepared! (course prerequisites)
  - CSCI 6212 Algorithms (or undergrad algorithms course)
  - An undergraduate operating systems course
- Be involved!
  - “Raise hand”, ask questions, discuss, etc.
  - Asynchronous opportunities will be available
- Be ready to code!
  - You will need to use **Go**, **Python** for your assignments
  - Mostly group projects



# CLASSES

- 2.5 hours is a long time for lectures!
  - We will try to break it up - discussions, demos, live coding
  - Some lectures may end early, with additional asynchronous material
- We want to make the best course we can for you!

# PARTICIPATE!

- You must “participate” 2X per week:
  - Attend lecture or office hours
  - Post a question/comment/answer on BB/Slack (during or outside of class)
- Examples:
  - Attend both lecture and office hours = 2 points 😊
  - Attend office hours and ask a question = 2 points 😊
  - Post 3 questions = 2 points 😊
  - Only attend lecture = 1 point ☹

# RESOURCES

- Slack: (linked from website, join after class)
- GitHub for collecting assignments
- Blackboard for grades, class meetings, and office hours
- Visual Studio Code - recommended IDE
  - Live share plugin allows group collaboration / help in office hours
- Repl.it - simple online editor for quick programming exercises
  - You can login with GitHub credentials if you want to save copies

# SEMESTER OUTLINE

- Building Blocks
  - Introduction to Distributed System and Cloud
  - Scalable Execution: Processes, threads, VMs, containers, parallelism vs concurrency
  - Communication: RPC, Message Oriented, Stream Oriented
- **Principles** of Distributed Systems
  - Coordination: Synchronization, Consistency, and Consensus
  - Reliability: Replication and Fault Tolerance
  - Performance: Metrics and Modeling Large Scale Systems
- Distributed Systems in **Practice**
  - Grid Computing
  - Cloud Computing
  - Web, Mobile, and IoT

4 Go programming assignments

Midterm

Large group project



# History of Computers

## Timeline and Ordering Activities

### INTRODUCTION

- Computer systems are undergoing revolution.
- Two advances in technology changed the game
  - 8bit → 16bit → 32bit → 64bit microprocessors
    - From a machine that cost \$10M and executed 1 inst./sec we have come to machine that cost \$1000 and execute 1 billion inst./sec
  - Computer networks LAN/WAN
    - From 64 Kbit/sec to Gigabit/Sec



# INTRODUCTION

- If we had this progress and improvement in cars industries:
  - A Rolls Royce would cost 1 dollar and get a billion miles per gallon.





WHAT IS THE CLOUD?



# WHAT IS THE CLOUD

- Giant warehouses
- 10s of thousands of servers
- Petabytes of storage
- 10s of thousands of Processor cores
- ... Interconnected...



# WHY INFRASTRUCTURE?

- Why do we need this amount of infrastructures?
  - Encyclopedia Britannica
    - - 40,000+ articles
    - - 32 hard bound volumes (32,640 pages)
  - Wikipedia
    - - 5,512,202 articles (in English)
    - - More than **5 TB** of text (about 7,500 CDs)
    - - More than 2000 volumes





# AND THEN BIG DATA

- Why do we need this amount of infrastructures?
  - Airbus A350
    - Contains around 6000 sensors across the entire plane that generates 2.5TB Data per day
  - Airbus A380-100
    - Expected to take the skies in 2020
    - Contains 10000 sensors just in each wings
  - Facebook
    - 20 TB photos each week
  - Google
    - 20000TB Data processing per day in 2008



# AND THEN BIG DATA

- Google Search Statistics

The average figure of how many people use Google a day, which translates into at least 2 trillion searches per year, 3.8 million searches per minute, 228 million searches per hour, and 5.6 billion searches per day.


- How much data do we generate?

According to the Forbes statistics:

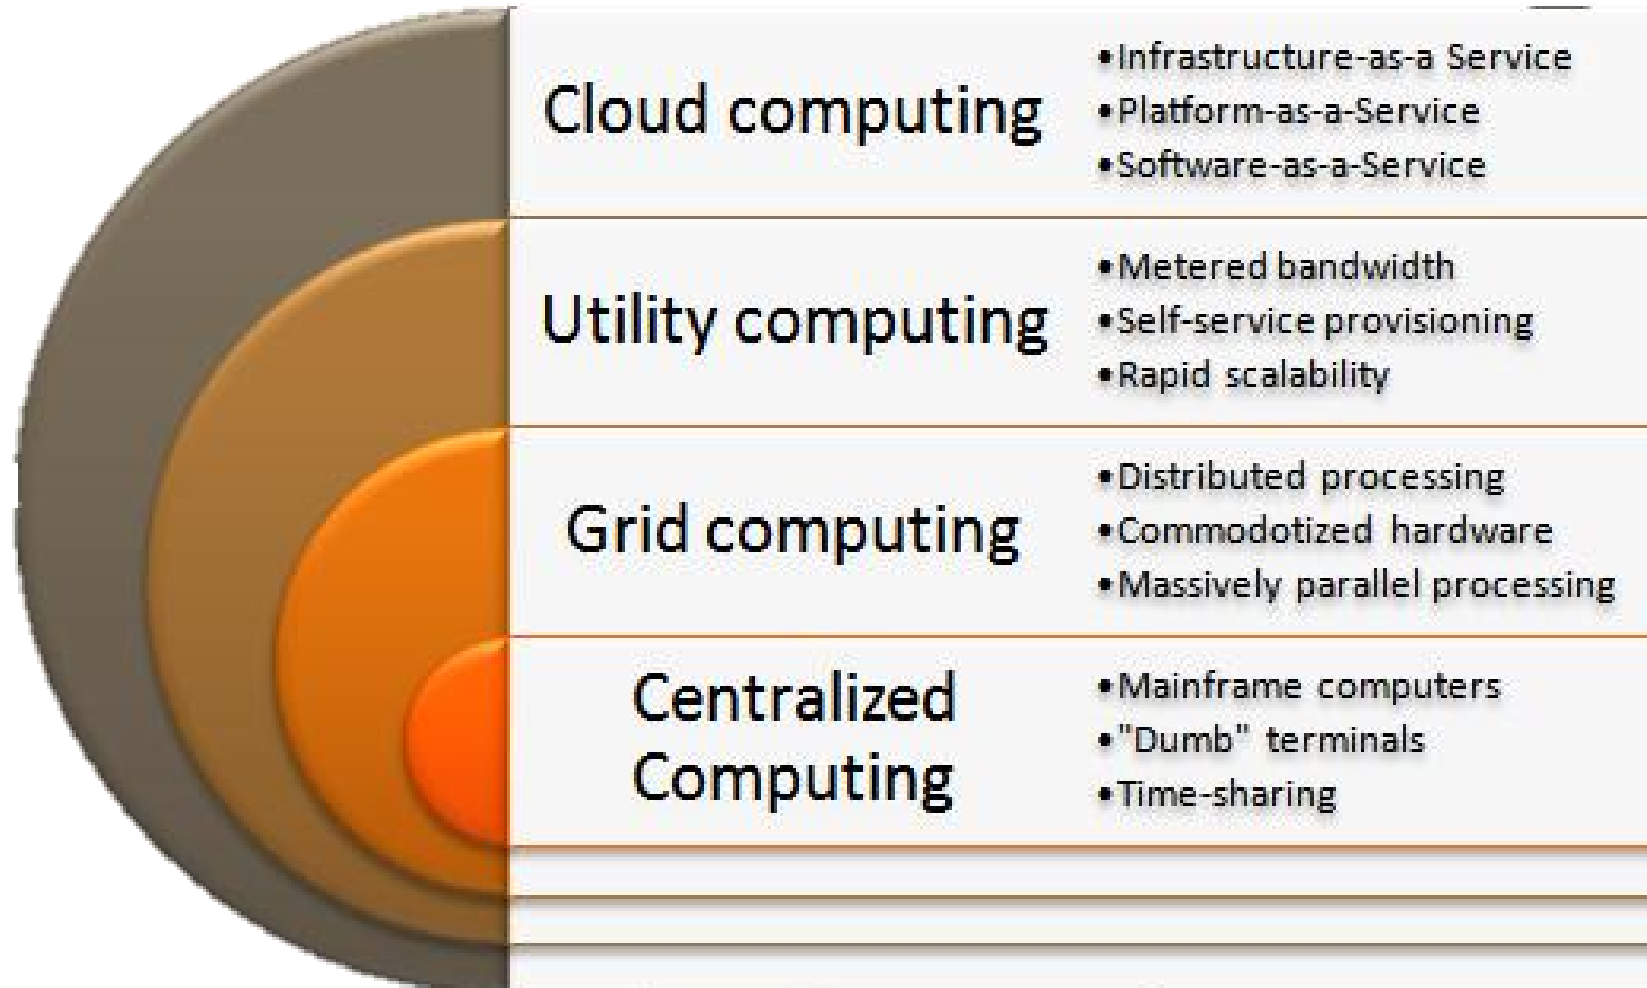
- 2.5 quintillion bytes of data created each day
- Over the last two years alone 90 percent of the data in the world was generated.

<b>KB</b>	<b>Kilo Byte</b>	<b>1 thousand bytes</b>
<b>MB</b>	<b>Mega Byte</b>	<b>1 million bytes</b>
<b>GB</b>	<b>Giga Byte</b>	<b>1 billion bytes</b>
<b>TB</b>	<b>Tera Byte</b>	<b>1 trillion bytes</b>
<b>PB</b>	<b>Peta Byte</b>	<b>1 quadrillion bytes</b>
<b>EB</b>	<b>Exa Byte</b>	<b>1 quintillion bytes</b>

# HISTORY OF CLOUD COMPUTING



# HISTORY OF CLOUD COMPUTING





# WHAT' S NEW

- There are four new features in the new generation of distributed and cloud systems:
  - Massive Scale
  - On-Demand Access: Pay-as-you-go
  - Data Intensive Nature: MBs became PBs and XBs
  - New Cloud Programming Paradigms: Map/Reduce Hadoop, Unstructured Data



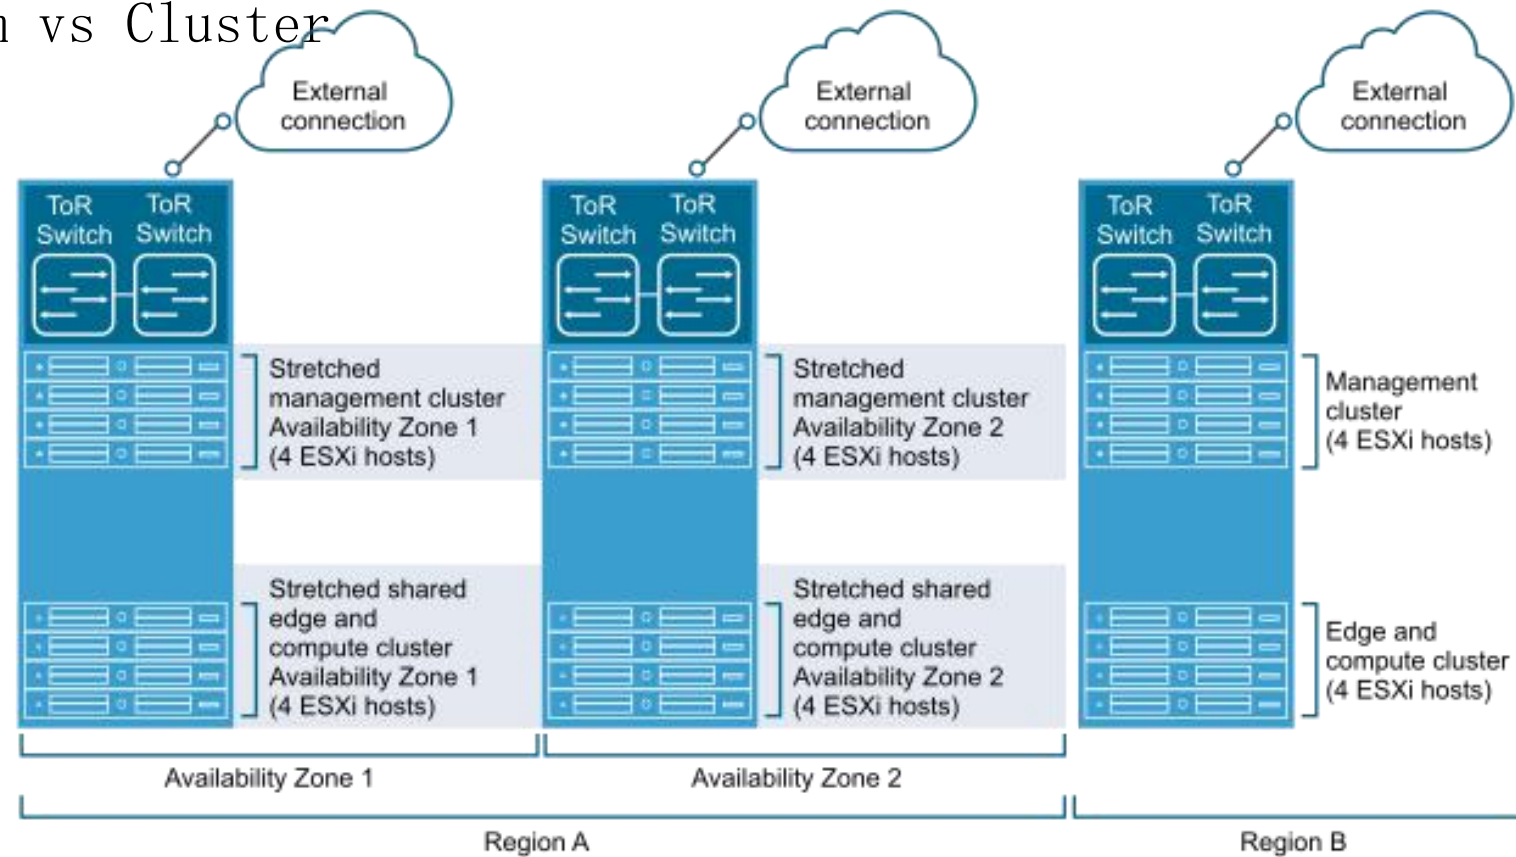
# \*<sub>aaS</sub> CLASSIFICATION

- HaaS : Hardware as a Service  
Hardware and backbone
- IaaS: Infrastructure as a Service  
AWS, Azure, GCP
- PaaS: Platform as a Service  
Google App engine, AWS Elastic Beanstalk
- SaaS: Software as a Service  
Google Doc, Dropbox

# CLOUD IS A ...

- Cloud vs Distributed System vs Cluster

- Client Server Architecture





# CLOUD IS A ...

- Can we say “ Cloud is a fancy word for a Distributed System?”

# WHAT IS A DISTRIBUTED SYSTEM

- A distributed system is a collection of independent computers that appears to its users as a single coherent system. [Andrew Tanenbaum]
  - distributed system consists of components that are autonomous
  - users (be they people or programs) think they are dealing with a single system. (Transparency)
  - distributed systems should also be relatively easy to expand or scale.
  - Heterogeneity
  - Concurrency



# GOALS OF DS

- Making resources accessible
- Distribution Transparency
  - Access
  - Location
  - Migration
  - Relocation
  - Replication
  - Concurrency
  - Failure
- Openness
- Scalability





# ACCESSIBILITY

- The main goal of a distributed system is to make it easy for the users and applications to access remote resources and to share them in a controlled and efficient way



# TRANSPARENCY

- **Transparency** in simple words is defined as the concealment from the user and the application programmer of the separation of components in a **distributed system**, so that the **system** is perceived as a whole rather than as a collection of independent components.



# OPENNESS

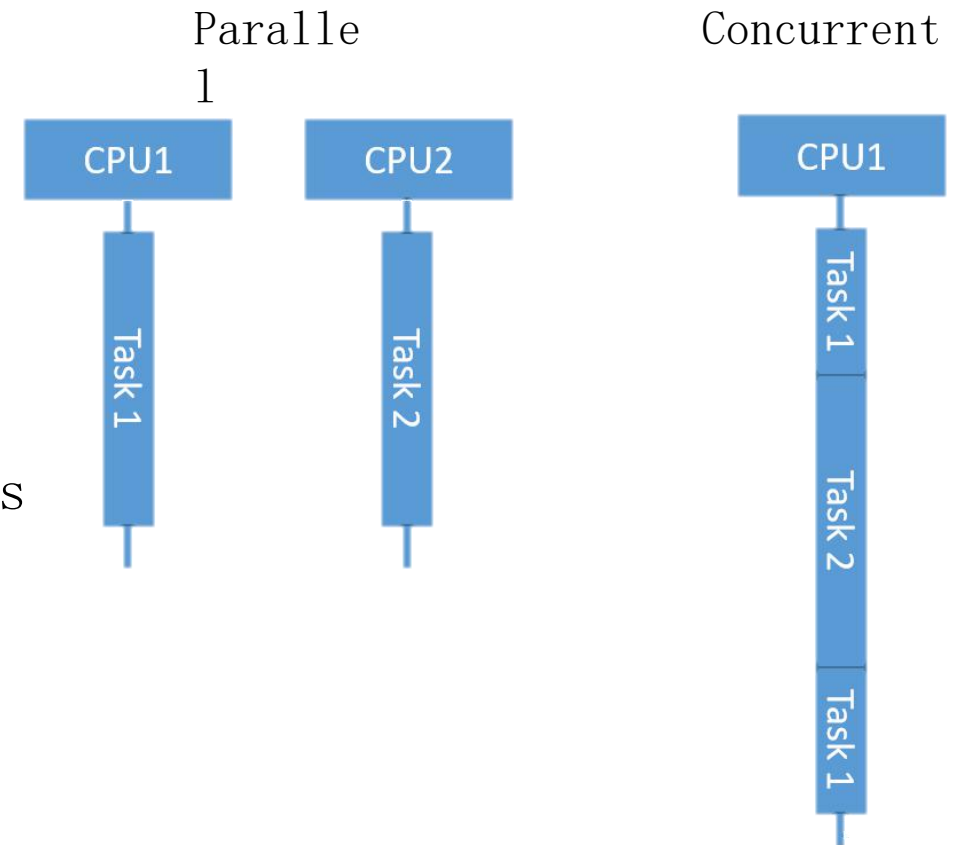
- An open distributed system is a system that offers services according to standard rules that describe the syntax and semantics of those services.

# SCALABILITY

- **Scalability** means you can increase or reduce the capacity, power or abilities of your system. It can be measured along at least three different dimensions:
  - A system can be scalable with respect to its size (add more users/resources to the system - can be consider as Scale up)
  - A geographically scalable system is one in which the users may lie far apart (Scale out)
  - A system can be administratively scalable. It means that it can still be easy to manage even if it spans many independent administrative organizations.

# CONCURRENCY VS PARALLELISM

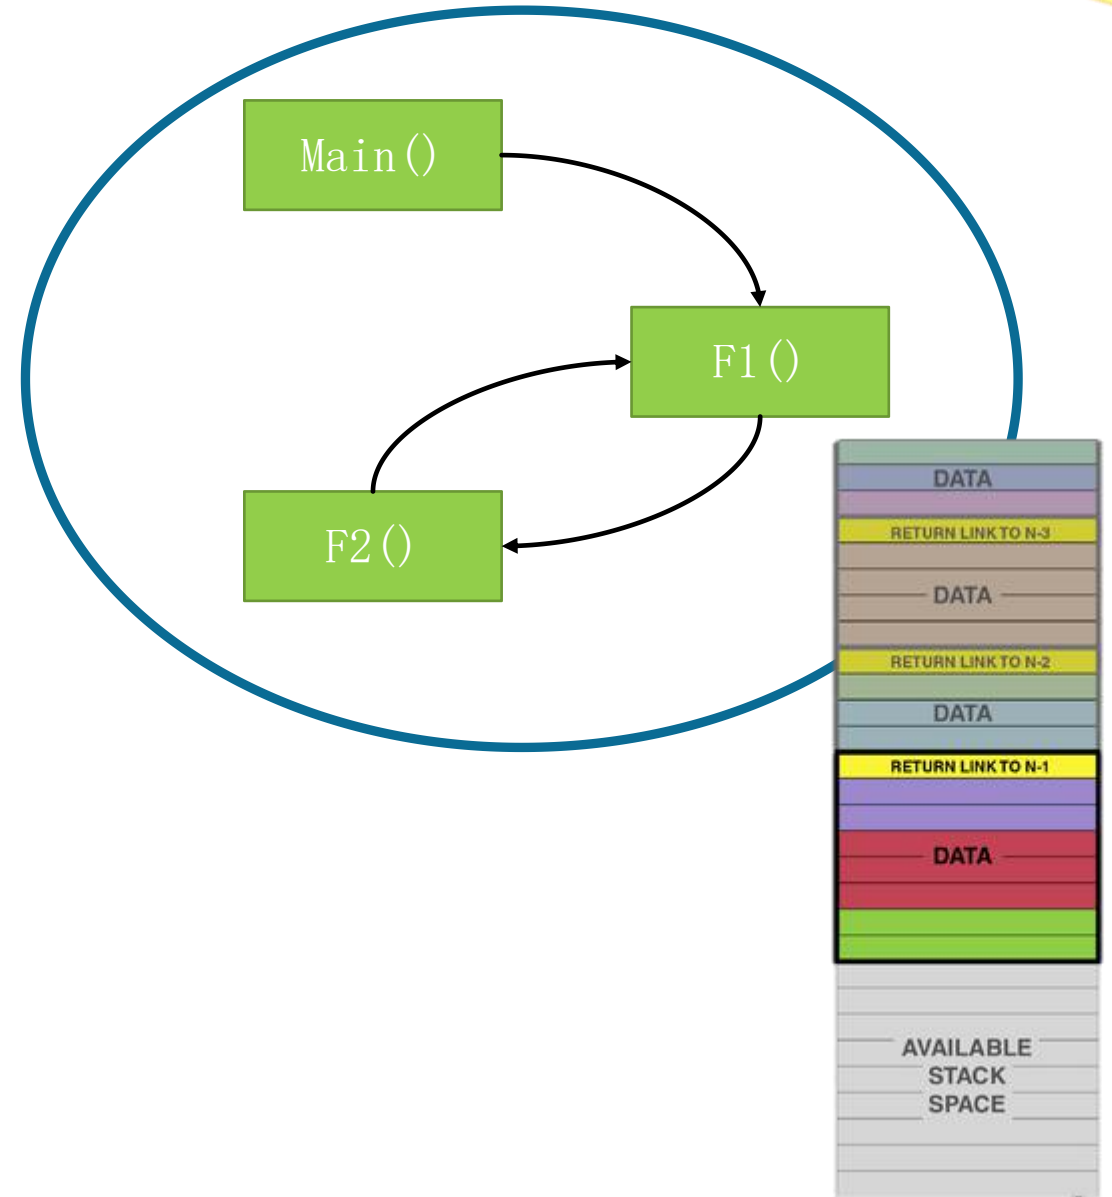
- Concurrency considers the checkpoints
- Parallelism considers time of progresses






# PROCESS

- Process
- Stack
- Program Counter
- Heap
- Etc.



# DISTRIBUTED ...

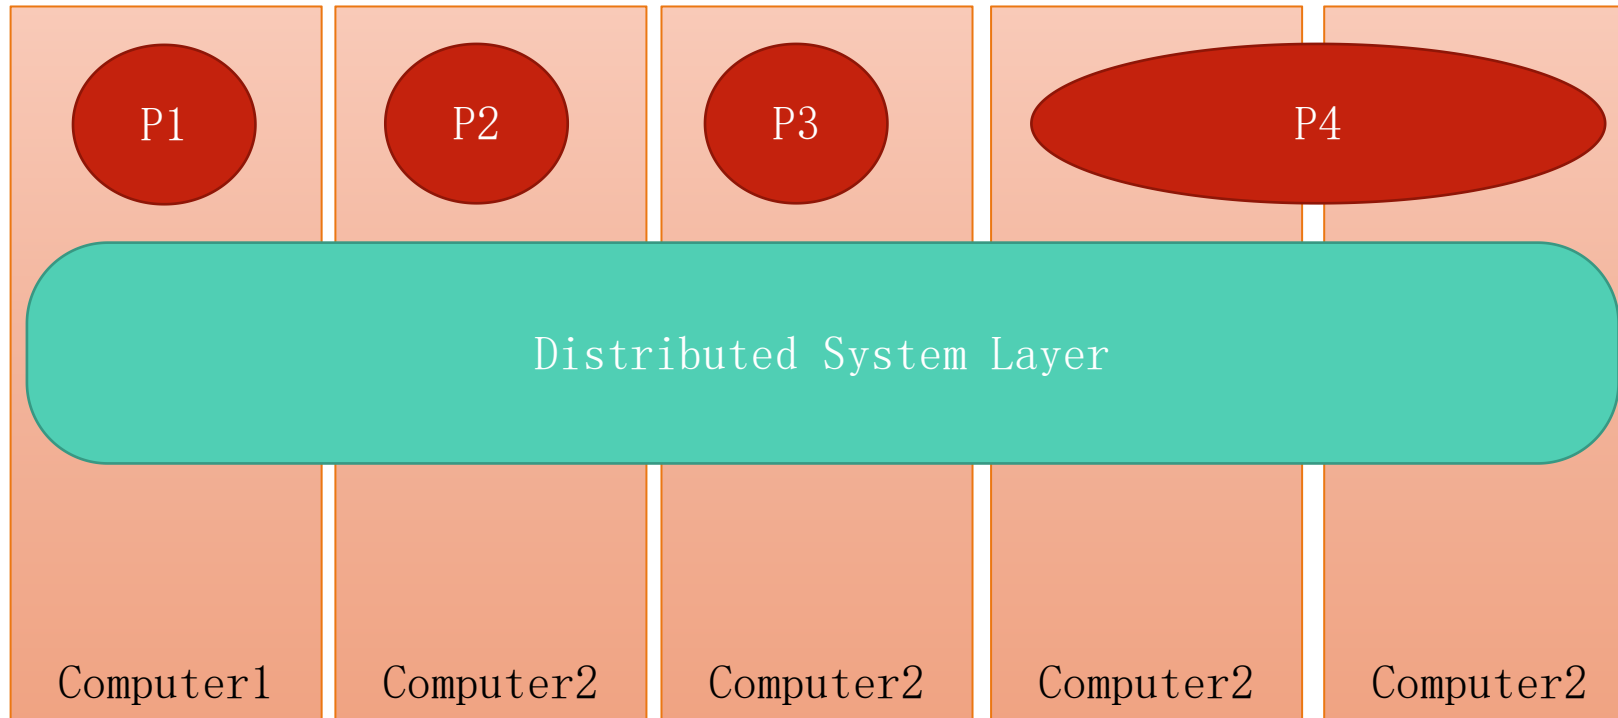
- Distributed System = Many Processes ?????



Reliable or Unreliable Communication

# HOW CAN WE HANDLE?

- Faster Computer Or Add Another Computer?



# BREAKOUT

- Grouping
  - ~5 people each
- 1. Introduce yourselves:
  - Who are you? What do you want to learn from this course?
- 2. Answer these questions as a group:
  - What is something you learned from the lecture so far?
  - What is a part of the lecture was confusing to you?
- Back to normal lecture in ~6 minutes!



# HW 1 : GO PARALLEL SUM



# PARALLEL SUM

- Assignment Goals:
  - Learn the basics of the Go programming language
  - Familiarize yourself with the editing environment and Git
  - Build two types of distributed systems
- This is an **individual** assignment
  - You must write all your own code
  - You may discuss general ideas with other students and link them help documentation
  - You may give general advice for debugging and design, but you should **never** have your code open while looking at someone else' s code!
  - This is more lenient than many classes, don' t abuse it!

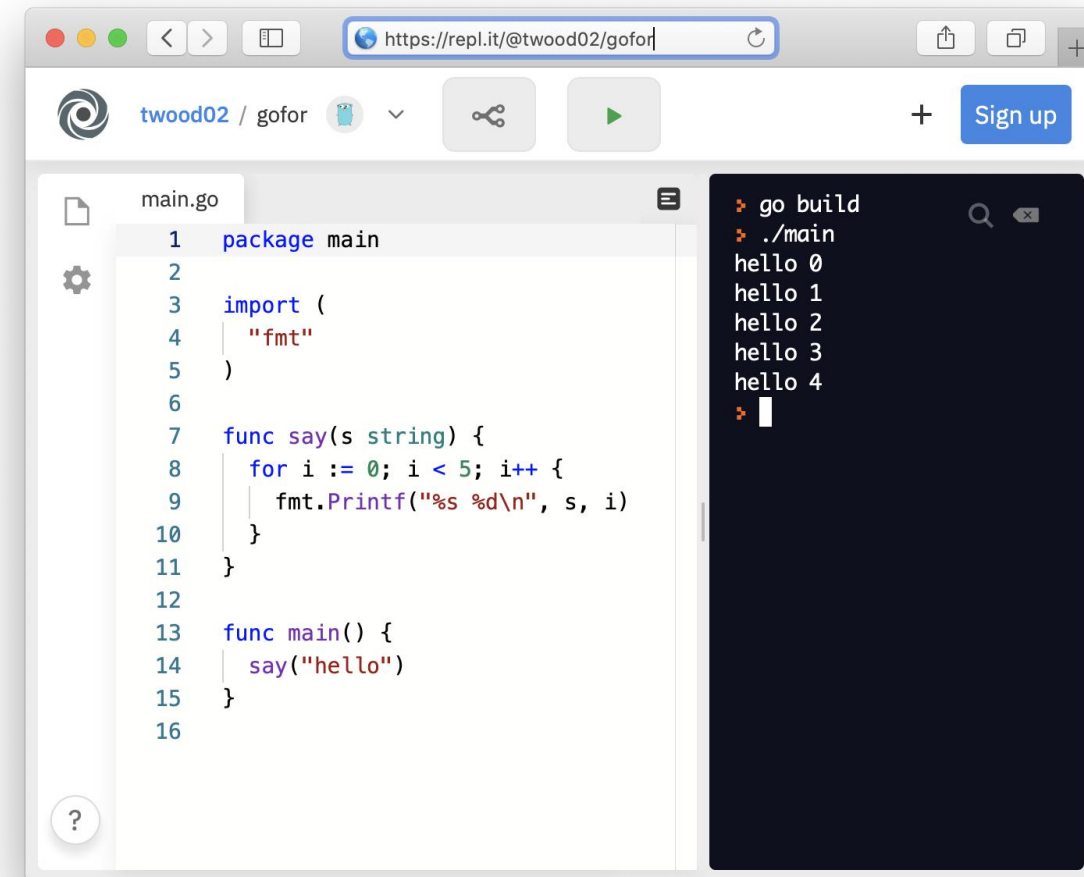
# WHY Go?

- Go has become a very popular language for building distributed systems
- Born at Google by Robert Griesemer, Rob Pike and Ken Thompson (C/Unix)
- Power and performance of C, but with the convenience and safety of more modern languages
- Learn more: <https://golang.org/doc/faq>

“Go ... [attempted] to combine the ease of programming of an interpreted, dynamically typed language with the efficiency and safety of a statically typed, compiled language. It also aimed to be modern, with support for networked and multicore computing.”

# PHASE 1: SEQUENTIAL SUM

- Starter code:
  - Reads a file and puts numbers in an array
- Your code:
  - Use a for loop and add up the numbers
  - Add command line parameter support
  - (this should be easy even if you've never touched go)
- Hint: Take a tour of Go
  - <https://tour.golang.org/list>



The screenshot shows a web-based Go Playground interface. The browser address bar displays `https://repl.it/@twood02/gofor`. The interface includes a user profile section for 'twood02 / gofor' with a share icon and a 'Sign up' button. The main area is divided into two panes. The left pane, titled 'main.go', contains the following Go code:

```
1 package main
2
3 import (
4     "fmt"
5 )
6
7 func say(s string) {
8     for i := 0; i < 5; i++ {
9         fmt.Printf("%s %d\n", s, i)
10    }
11 }
12
13 func main() {
14     say("hello")
15 }
16
```

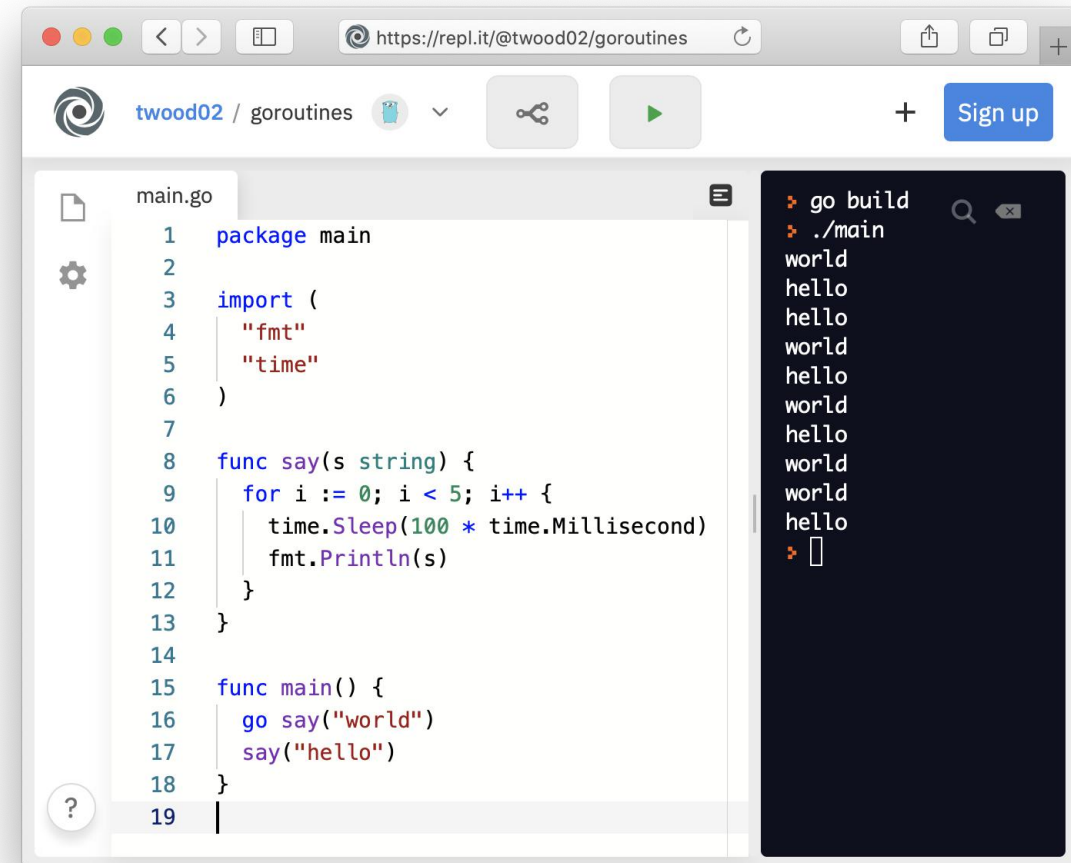
The right pane shows the output of the program after running the command `go build` and `./main`. The output is:

```
hello 0
hello 1
hello 2
hello 3
hello 4
```

<https://repl.it/@twood02/gofor>

# PHASE 2: PARALLEL SUM

- Main thread still reads in file and makes array (see starter code)
- Use Goroutines to parallelize the addition
  - A **Goroutine** is a lightweight thread
  - What does this mean with regards to concurrency and parallelism?
- How will the main thread and goroutines coordinate?
  - Need to pass numbers to be summed
  - Need to get back the result
  - Hint: learn about **Go Channels**!



The screenshot shows a web-based Go IDE interface. The main editor displays a Go program named `main.go` with the following code:

```
1 package main
2
3 import (
4     "fmt"
5     "time"
6 )
7
8 func say(s string) {
9     for i := 0; i < 5; i++ {
10         time.Sleep(100 * time.Millisecond)
11         fmt.Println(s)
12     }
13 }
14
15 func main() {
16     go say("world")
17     say("hello")
18 }
19
```

On the right side, the terminal output shows the result of running the program:

```
go build
./main
world
hello
hello
world
hello
world
hello
world
hello
world
hello
```

<https://repl.it/@twood02/goroutines>

# PHASE 3: HTTP+RPC

- Let's make a “real” distributed system! Two Go programs:
- HTTP Frontend
  - Accepts a client request specifying file to process
- RPC Backend
  - Receives a Remote Procedure Call from frontend to trigger the summation
  - Uses goroutines to parallelize like in prior phase

